

Modern key-value stores use LSM-based storage, where unlike traditional index structures, these do not perform in-place updates. Rather, LSM tree first buffers all writes in main memory, and subsequently flushes the buffer as sorted run to disk whenever it fills up, and organizing the disk-runs into a number of levels of increasing sizes. LSM-tree later sort-merges these runs. This design has lots of benefits including superior write performance, high space utilization, tunability, and simplification of concurrency control and recovery. Two well known optimizations used by most LSM-trees include Bloom filters (a probabilistic data-structure designed to improve point lookup costs) and Partitioning. Recent LSM-trees typically organize their memory components using skip-list or B+ tree and organize their disk components using B+ trees or sorted-string tables (SSTables).

In this paper, the authors survey the recent research efforts on improving LSM-trees. The six major issues and some of their optimizations discussed by the authors are described below.

Write Amplification, where it impacts both the write performance of LSM-tree and also the lifespan of SSDs due to frequent disk writes. Some possible improvements provided in the paper include Tiering, which has lower write amplification than levelling. Authors describe four structures which are variants of partitioned tiering with vertical grouping such as: WriteBuffer(WB) which uses hash-partitioning to achieve workload balance, light-weight compaction tree (LWC-tree) where if a SSTable group contains too many entries, it will shrink the key range after it is merged, PebblesDB (built using Fragmented Log-Structured Merge Tree, uses the idea of guards) and dCompaction which uses the idea of virtual SSTables and virtual merges to reduce merge frequency. Merge skipping is also discussed as a way to improve write performance where main idea is, if possible, to push some entries directly from level 0 to a higher level by skipping some level-by-level merges then total write cost is reduced. Data skew is exploited where basic idea is to separate hot keys from cold keys in the memory so that only cold keys are flushed to the disk.

Merge Operations is the second issue where they can have negative impacts on the system including buffer cache misses after merges and write stalls during large merges. Some optimizations mentioned include improving merge performance, where VT-tree is described which presents a stitching operation to improve performance. Pipelined merge is also discussed which uses CPU and I/O parallelism. Reducing Buffer Cache Misses is discussed where an approach is to use a smart cache warmup algorithm to fetch new component incrementally, to smoothly redirect incoming queries from old components to new components. Another optimization discussed is Minimizing Writing Stalls, where bLSM proposes a spring-and-gear merge scheduler to minimize write stalls, where basic idea is to tolerate an extra component at each level so that merges at different levels can proceed parallelly.

Hardware is the third issue discussed, where LSM-tree implementation must be such as to fully utilize the underlying platforms. Previously hard disks were the main hardware platforms to be used, but recent underlying platforms include large memory which improves both write and lookup performances, multi-core which includes an optimistic concurrency control approach, SSDs which support efficient random I/Os (unlike traditional hard disks which only supported sequential I/Os) and NVMs which further provide efficient byte-addressable random accesses with persistence guarantees, and attempts to perform native management of storage devices (SSDs/HDDs) to optimize the performance of LSM-tree implementations.

Handling special workloads is the fourth issue, where certain special workloads such as temporal data, small data, semi-sorted data and append-mostly data can also be considered to achieve better performance. Basic LSM-tree structure is adapted to handle special workloads. LHAM more efficiently supports temporal workloads by attaching a range of timestamps to each component to facilitate temporal queries processing by pruning irrelevant components. LSM-trie manages a large number of key-value pairs where each key-value pair is small, proposing a number of optimizations to reduce meta-data overhead. SlimDB is discussed which targets semi-sorted data by adopting a hybrid structure with tiering on lower levels and leveling on higher ones, also uses multi-level cuckoo

filters to improve point lookup performance. Two newly proposed merge policies were discussed which optimized for append-mostly workloads with a bounded number of components (MinLatency, Binomial).

Auto-tuning is the fifth issue where tunability along the pareto-curve is to be determined since a method can't be write, read and space-optimal at the same time. Some optimizations discussed are parameter tuning, Monkey being an example which co-tunes the merge policy, size ratio and memory allocation between main memory and Bloom filters to find optimal LSM-tree design for a given workload. Tuning Merge policies is discussed, an example being the ChooseBest policy which selects an SSTable with the fewest overlapping SSTables at the next level to merge. Dynamic Bloom filter memory allocation is also discussed in this paper, an example being ElasticBF which dynamically adjusts bloom filter false positive rates based on data hotness and access frequency to optimize reads.

Secondary indexing is the final issue discussed, which is needed to support efficient queries on non-key attributes. Index structures (such as LSII), index maintenance, statistics collection, distributed indexing all focus on improving LSM-trees in database settings with secondary indexes and other auxiliary structures.

In conclusion, this paper surveyed the recent research efforts done to improve LSM-trees. General aspects of improvements were described and the optimizations done in each of them were discussed in great detail. Furthermore, it provides a review of five representative LSM-based opensource NoSQL systems such as LevelDB, RocksDB, Cassandra, HBase, and AsterixDB. Some future works include thorough performance evaluation, where comparisons are to be done against well-tuned baseline LSM-tree (LevelDB or RocksDB) rather than untuned configuration. Furthermore, performance evaluation for space utilization was neglected, so, for future works, this may be considered as a significant factor. Partitioned Tiering structure may be further examined whereby horizontal and vertical grouping can be used to possibly design a new scheme that has the advantages of both. Hybrid Merge Policy may be studied which takes the advantages of both levelling and tiering. Performance Variance of LSM-trees is as important metric as absolute throughput, so a future work may be to design mechanisms to minimize this.